

Kubernetes (k8s)

Rajesh G

CTO, Managing Partner

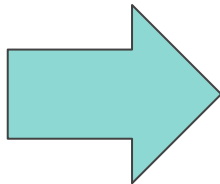
<https://unigps.in>





Training Objectives

At the end of training,
participants should be able to



- ❑ Know Kubernetes and Be a Helmsman
- ❑ Create and run PODs
- ❑ Bundle applications & Deploy
- ❑ Service apps using Load Balancers
- ❑ Troubleshoot

Table of Contents

DAY ONE			
Docker	Kubernetes Core	Pod	Multi Container Pods
Dockerfile	Architecture	Lifecycle	Ambassador
Images & Containers	Components	Create & Launch Pods	Side car
Registry	Kubernetes API Primitives	Demos	Adapter
Jenkins / CI for images	Kubectrl	Practicals	Demos
Demos	Demos		Practicals
Practicals	Practicals		
DAY TWO			
Pod Design	Persistence	Configuration	Scheduling
Deployments	Persistence Volume	ConfigMaps	NodeName
Rolling Updates	Persistence VolumeClaim	SecurityContexts	Taints
Rollbacks	Stateless	Resource Requirements	Tolerations
Jobs, Cron Jobs	Statefull	Secrets	Affinity
Labels, Selectors, Annotations	Demos	Service Accounts	Demos
Demos	Practicals	Demos	Practicals
Practicals		Practicals	
DAY THREE			
Services	Ingress & Networking	Monitoring	Recap
ClusterIP	Ingress Controller	LivenessProbes	Review
NodePort	Ingress Resources	ReadinessProbes	Q&A
LoadBalancer	Network Policies	Container Logging	
Connecting applications with Services	Host mapping with service	Monitoring & Debugging	
App deployment - Multi node cluster	TLS	Demos	
Auto Scaling pods	Multi-domain frontends with service	Practicals	
Auto Scaling nodes in cluster	Demos		
Demos	Practicals		
Practicals			



Docker

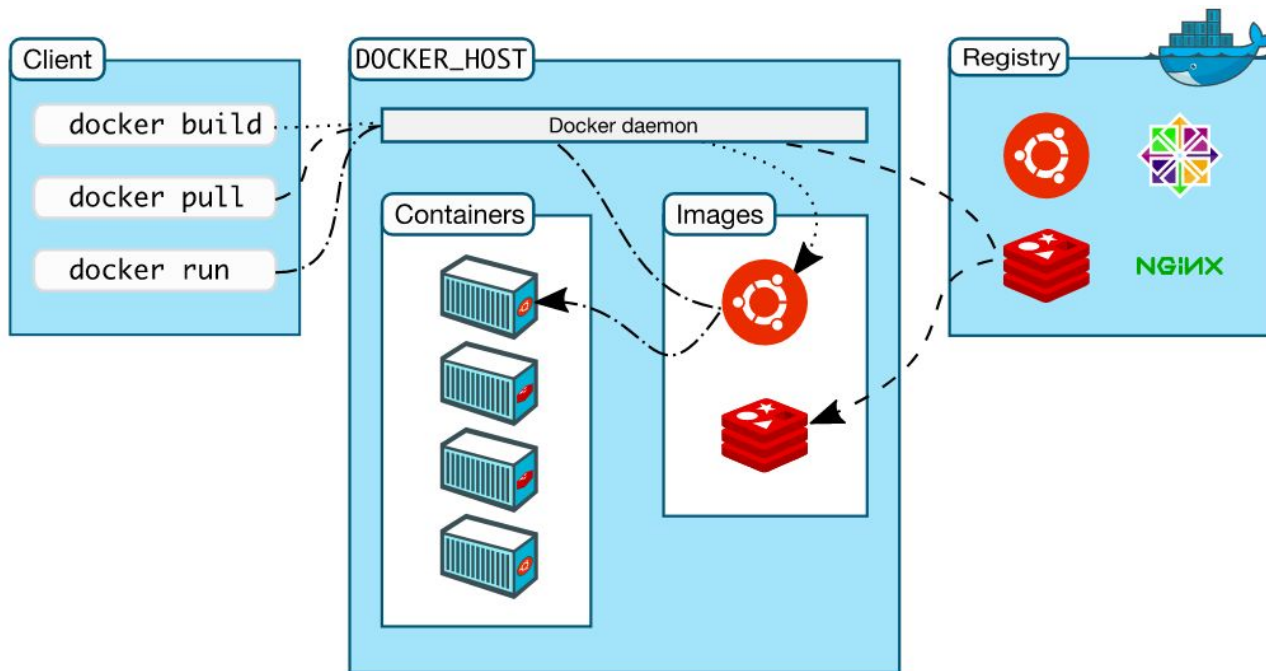
- Overview
- Dockerfile
- Images & Containers
- Registry
- Jenkins / CI
- Demo



Containers - Benefits


- Flexible: Even the most complex applications can be containerized.
- Lightweight: Containers leverage and share the host kernel.
- Interchangeable: You can deploy updates and upgrades on-the-fly.
- Portable: You can build locally, deploy to the cloud, and run anywhere.
- Scalable: You can increase and automatically distribute container replicas.
- Stackable: You can stack services vertically and on-the-fly
- Running more workload on the same hardware

Docker Architecture





Docker Architecture





Introducing the Dockerfile

A Dockerfile is a text document that contains

- a set of instructions required to assemble the app (image) and/ run it

Usage:

```
docker build [OPTIONS] PATH | URL | -
```

Options:

<code>--add-host list</code>	Add a custom host-to-IP mapping (host:ip)
<code>--compress</code>	Compress the build context using gzip
<code>--cpu-quota int</code>	Limit the CPU CFS (Completely Fair Scheduler) quota
<code>-f, --file string</code>	Name of the Dockerfile (Default is 'PATH/Dockerfile')
<code>--force-rm</code>	Always remove intermediate containers
<code>--label list</code>	Set metadata for an image
<code>-m, --memory bytes</code>	Memory limit
<code>--pull</code>	Always attempt to pull a newer version of the image
<code>--rm</code>	Remove intermediate containers after a successful build (default true)
<code>-t, --tag list</code>	Name and optionally a tag in the 'name:tag' format



Introducing the Dockerfile

Example:

- `docker build -f Dockerfile .`

```
rajesh@rajesh-Gazelle:~/git/dockers/trainora-aug/images/simple$ cat Dockerfile
```

```
FROM alpine:latest
```

```
MAINTAINER rajesh@unigps.in
```



Introducing the Dockerfile

Few more variations:

- `docker build -t myfirstimage -f Dockerfile .`
- `docker build -f /home/rajesh/git/dockers/trainora-aug/images/simple/Dockerfile-myfirstimage .`
- `docker build -t myfirstimage -f ./simple/Dockerfile ./simple/`
- `docker build -t myimage -t rajesh/myimage:1.0.0 -t localhost:5000/rajesh/myimage:1.0.0 .`



Introducing the Dockerfile

- ENV - to set environment variables
- EXPOSE - to expose ports
- FROM - base image
- LABEL - to add metadata to image
- HEALTHCHECK - to check if container is running
- USER - to set user and group
- VOLUME - to specify mount point from external host
- WORKDIR - workdir to run any of the commands



Introducing the Dockerfile

- ARG - variable used during build time
- CMD - to provide defaults to executing container
- RUN - to execute commands in new layer
- COPY - Copy file,dir or remote url to image
- ADD - Copy file,dir or remote url to image
- ENTRYPOINT - to configure container as executable
- MAINTAINER - the image maintainer

RUN COPY ADD instructions create new layers in the image stack - refer layering section



Building Images (Python)

rajesh@rajesh-Gazelle:~/git/dockers/trainora-aug/images/python\$ cat Dockerfile

```
FROM python:2.7-slim
WORKDIR /app
ADD app.py /app
ADD requirements.txt /app
RUN pip install --trusted-host pypi.python.org -r requirements.txt
EXPOSE 80
ENV name world
CMD ["python", "app.py"]
```

Build

- `docker build -t mypython .`

Run

- `docker run -p 80:80 mypython`



Dockerfile - Example (Apache)

```
FROM bitnami/minideb-extras:jessie-r23
LABEL maintainer "Bitnami <containers@bitnami.com>"

# Install required system packages and dependencies
RUN install_packages libapr1 libaprutil1 libc6 libexpat1 libffi6 libgmp10 libgnutls-deb0-28 libhogweed2 libldap-2.4-2 libnettle4
libp11-kit0 libpcre3 libsasl2-2 libssl1.0.0 libtasn1-6 libuuid1 zlib1g
RUN bitnami-pkg unpack apache-2.4.29-1 --checksum
42114e87aafb1d519ab33451b6836873bca125d78ce7423c5f7f1de4a7198596
RUN ln -sf /opt/bitnami/apache/htdocs /app

COPY rootfs /

ENV APACHE_HTTPS_PORT_NUMBER="443" \
    APACHE_HTTP_PORT_NUMBER="80" \
    BITNAMI_APP_NAME="apache" \
    BITNAMI_IMAGE_VERSION="2.4.29-r1" \
    PATH="/opt/bitnami/apache/bin:$PATH"

EXPOSE 80 443

WORKDIR /app
ENTRYPOINT ["/app-entrypoint.sh"]
CMD ["nami", "start", "--foreground", "apache"]
```



Dockerfile - Example (Jenkins CI)

```
FROM jenkinsci/jenkins:latest
LABEL maintainer "r1co@post-box.cc"

USER root

# install docker cli
RUN mkdir -p /tmp/_install && cd /tmp/_install && wget https://get.docker.com/builds/Linux/x86_64/docker-latest.tgz && tar -xvzf
docker-latest.tgz && cd docker && cp docker /usr/bin/docker && rm -rf /tmp/_install
RUN chmod +x /usr/bin/docker
# add jenkins to docker group
RUN groupadd -g 999 docker
RUN usermod -a -G docker jenkins
# install docker-compose
RUN curl -L https://github.com/docker/compose/releases/download/1.7.1/docker-compose-`uname -s`-`uname -m` >
/usr/local/bin/docker-compose
RUN chmod +x /usr/local/bin/docker-compose
USER jenkins
```



Dockerfile - Example (Multi stage)

```
FROM golang:1.7.3 AS builder
WORKDIR /go/src/github.com/alexellis/href-counter/
RUN go get -d -v golang.org/x/net/html
COPY app.go .
RUN CGO_ENABLED=0 GOOS=linux go build -a -installsuffix cgo -o app .

FROM alpine:latest
RUN apk --no-cache add ca-certificates
WORKDIR /root/
COPY --from=builder /go/src/github.com/alexellis/href-counter/app .
CMD ["/app"]
```




Docker Hub - store & retrieve

<https://hub.docker.com> (register and create login)

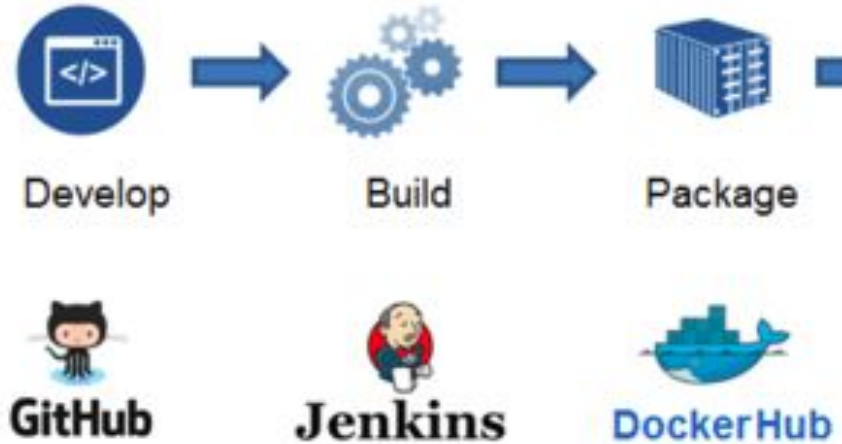
- `docker tag alpine rajeshgheware/alpine:rajesh`
- `docker push rajeshgheware/alpine:rajesh`
- `docker pull rajeshgheware/alpine:rajesh`



Registry Server

- With no docker volume (uses default volume for container)
 - `docker run -d -p 5000:5000 --name registry registry:2`
 - `docker push localhost:5000/rajesh/alpine:test`
 - `Docker pull localhost:5000/rajesh/alpine:test`
- With docker volume
 - `docker volume create docker_registry`
 - `docker run -d -p 5000:5000 -v docker_registry:/var/lib/registry --name registry registry:2`
 - `docker container stop registry && docker container rm -v registry`
- With Volume Mount on Host
 - `docker run -d -p 5000:5000 -v /media/deepti/Ubuntu/home/docker_registry:/var/lib/registry --name registry registry:2`

Build Image using CI / Jenkins



```
docker run -p 8080:8080 -p 50000:50000 -v /var/run/docker.sock:/var/run/docker.sock rlco/jenkins-docker
```

<https://jenkins.io/doc/tutorials/build-a-node-js-and-react-app-with-npm/> (ci for node reactjs build with pipeline)

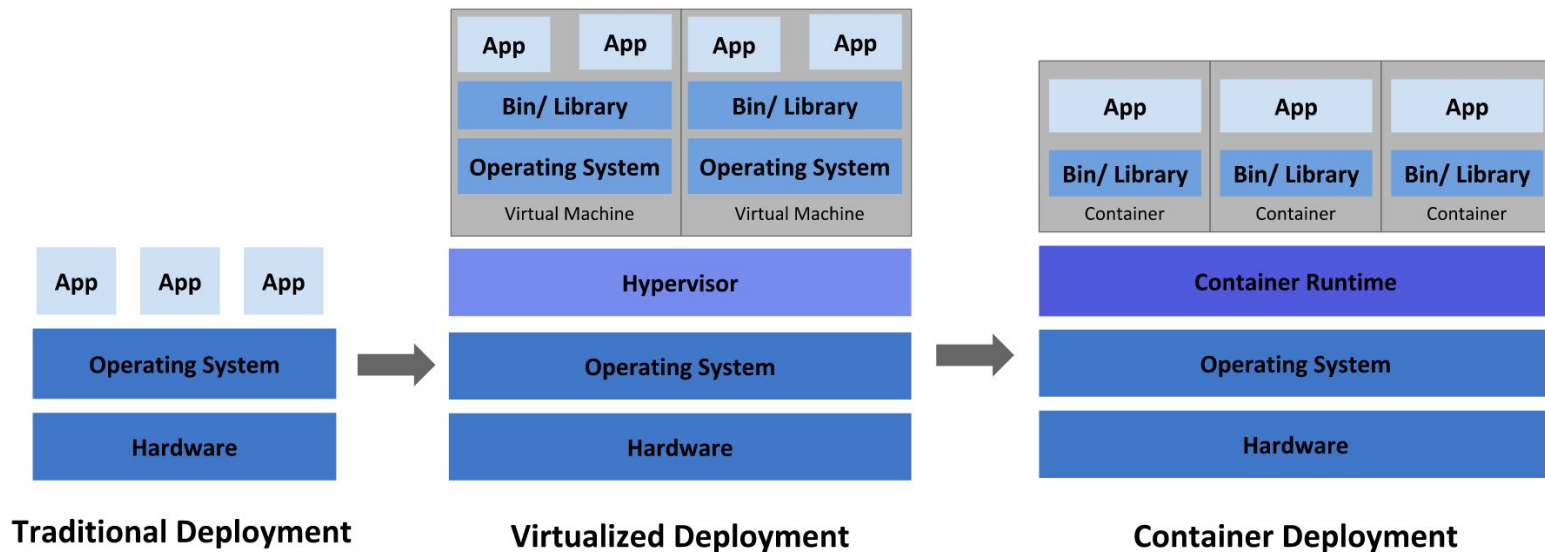


Kubernetes Core

- Architecture
- Components
 - Master Components
 - Node Components
 - Add ons
- API Primitives
- Kubectl
- Demo
- Practicals



Deployment - Journey





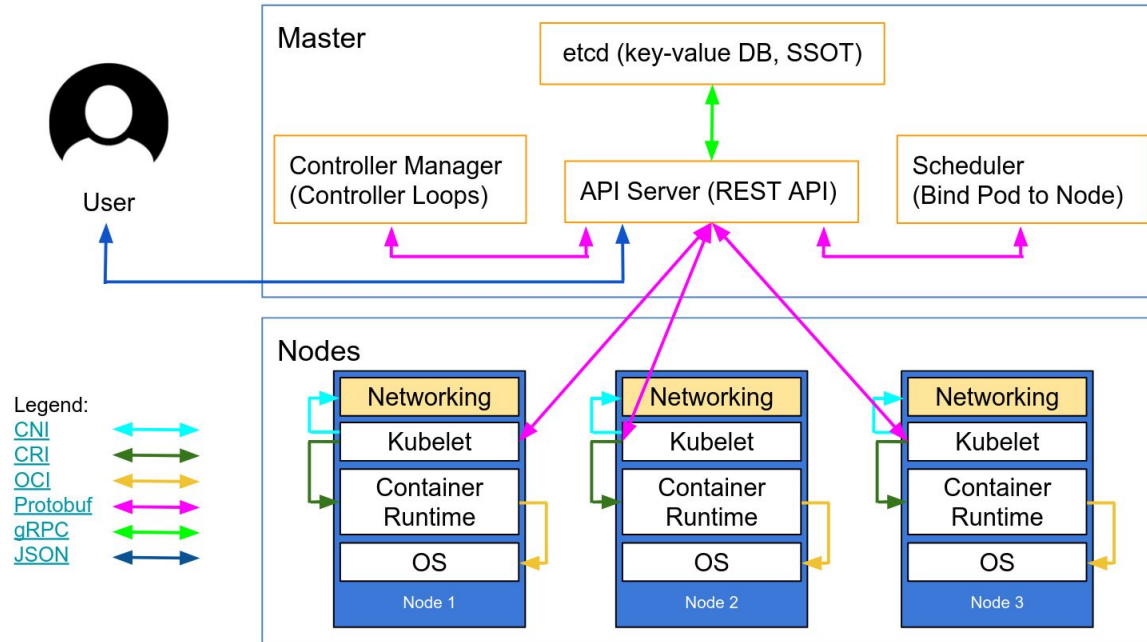
What is / why Kubernetes

Kubernetes - Helmsman (in ancient greek): Guy who steers ship / boat

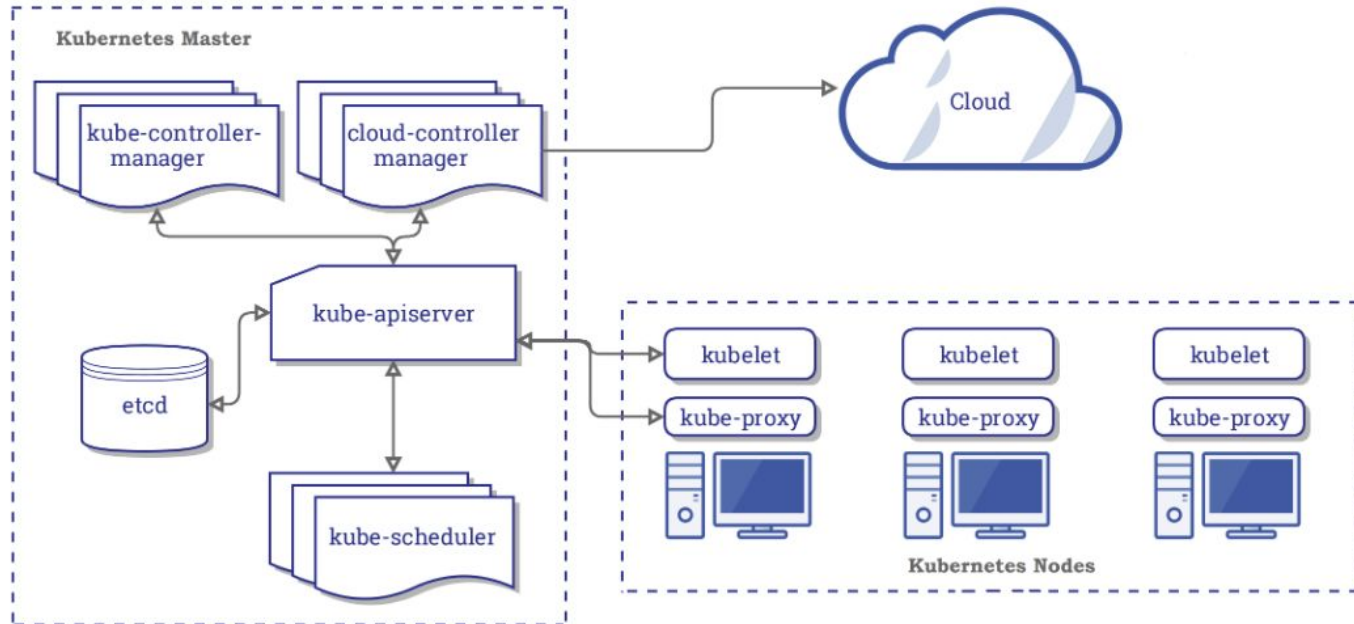
Why Kubernetes?

- Service Discovery & Load Balancing
- Storage Orchestration
- Automated rollouts & rollbacks
- Automatic bin packing
- Self-healing
- Secret and configuration Management

Architecture - Overview



Architecture (view 2)





Master Components - ETCD

- Distributed reliable key-value store that is simple, secure & fast
- Uses RAFT based consensus algorithm to work in distributed environment
- Key value store distributed database
- Runs on port 2379



Master Components - API Server

- The central management entity
- Only component that connects to ETCD
- Designed for horizontal scaling

Connectivity:

- External: kubectl
- Internal: kubelet
- Persistent Storage: ETCD



Master Components - Scheduler

Schedules pods on appropriate Node(s)

Watches for newly created PODs that have no nodes assigned

Decision Parameters:

- Resource requirements (memory, cpu, disk type say SSD)
- Hardware, Software, Policy requirements
- Affinity, Anti-affinity
- Data locality
- Inter workload interference
- Deadlines



Master Components - Kube Controller

- Node Controller
 - Responsible for noticing and responding when nodes go down
- Replication Controller
 - Responsible for maintaining the correct number of pods for every replication controller object in the system
- Endpoints Controller
 - Populates the Endpoints object (that is, joins Services & Pods)
- Service Account & token Controller
 - Create default accounts and API access tokens for new namespaces



Master Components - Cloud Controller

- Route Controller
 - For checking the cloud provider to determine if a node has been deleted in the cloud after it stops responding
- Service Controller
 - For setting up routes in the underlying cloud infrastructure
- Service Controller
 - For creating, updating and deleting cloud provider load balancers
- Volume Controller
 - For creating, attaching, and mounting volumes, and interacting with the cloud provider to orchestrate volumes



Node Components - kube-proxy

- Network proxy that runs on every node in cluster
- Maintains network rules on nodes
- Uses OS packet filtering layer else forwards traffic itself



Node Components - Container RT

- Docker
- Containerd
- cri-o
- rktlet

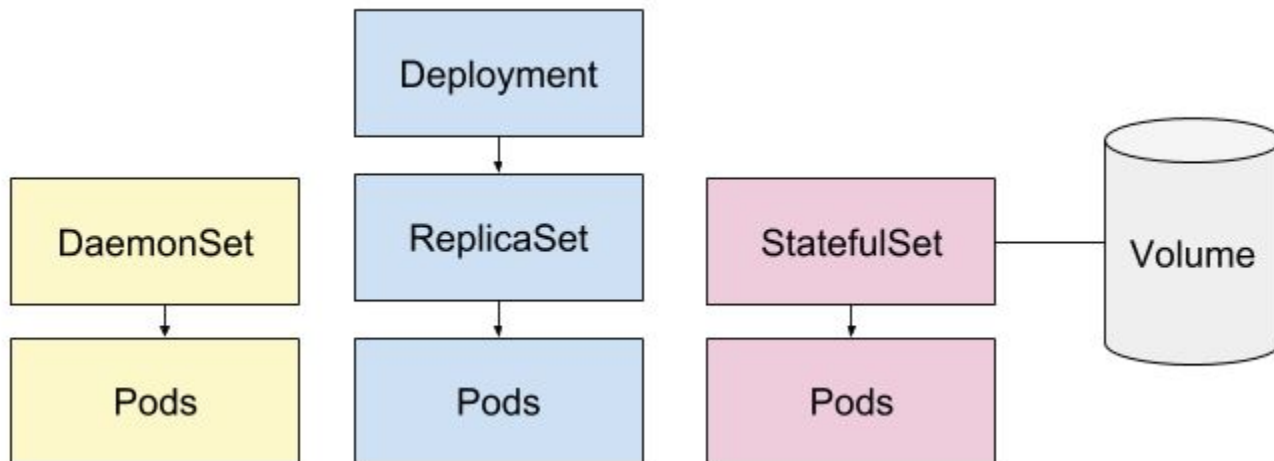


Addon Components

- Cluster DNS
 - Cluster DNS is a DNS server, in addition to the other DNS server(s) in your environment, which serves DNS records for Kubernetes services
- Web UI
 - General purpose, web-based UI for Kubernetes clusters to view and manager cluster
- Container Resource Monitoring
 - Generic time-series metrics about containers in a central database, and provides a UI for browsing that data
- Cluster level Logging
 - Mechanism responsible for saving container logs to a central log store with search/browsing interface

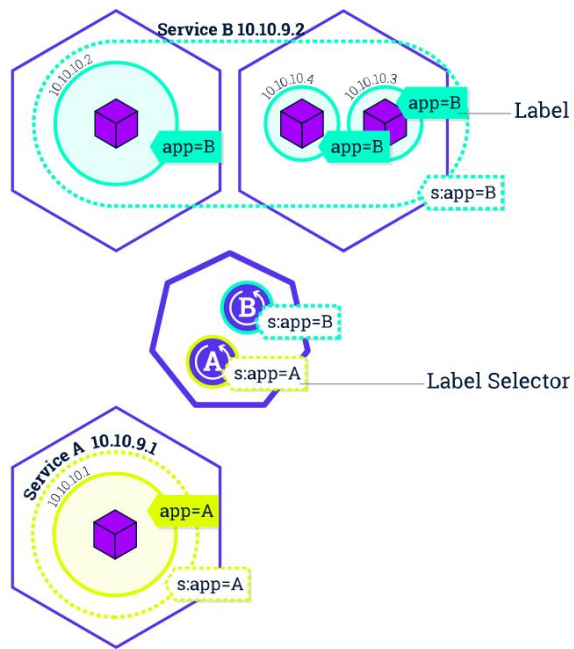


Objects





Objects





Node Components - kubelet

- Runs on every node
- Ensures containers are running & healthy in PODs
- Doesn't manage container not created by K8S



Kubectl

Command line tool to control kubernetes cluster

- Imperative commands to manage objects (basic & intermediate)
- Deploy commands
- Cluster Management commands
- Troubleshooting and Debugging
- Advanced, Settings and Other



kubectl - commands

- `kubectl get pods`
- `kubectl describe pod hello-world`
- `kubectl describe pod/nginx`
- `kubectl delete pod nginx`
- `kubectl cluster-info`
- `kubectl get pods -o yaml`
- `kubectl get services -o json`
- `kubectl get pods --sort-by=.metadata.name`
- `kubectl get rs,deployments,service`
- `kubectl describe pods`
- `kubectl get pod/<pod-name> svc/<svc-name>`
- `kubectl get pod -l name=<label-name>`
- `kubectl delete pods --all`
- `kubectl get nodes -o json | jq '.items[] | {name:.metadata.name, cap:.status.capacity}'`
- `kubectl get nodes -o yaml | egrep "\sname:|cpu:|memory:"`
- `kubectl get all`
- `kubectl run hello --image=tutum/hello-world --port=80`
- `kubectl run -it busybox --image=busybox --restart=Never`
- `kubectl run nginx --image=nginx`



Exercises (30 mins)

Run Hello World POD using tutum/hello-world image (kubectl run...) & then

- Get POD summary (kubectl get ...)
- Get POD details (kubectl describe ...)
- Get POD IP (kubectl describe pod... -o yaml | egrep....podIP:)
- Delete the POD created above (kubectl delete ...)
- Verify using kubectl get all
- View cluster info



POD

- Overview
- Lifecycle
- Init Containers
- Preset
- Topology Spread
- Ephemeral Containers



POD - Overview

- ❖ Smallest deployable unit
- ❖ Supports multiple cooperating processes (containers) that form cohesive unit of service
- ❖ Ephemeral Entity

Encapsulates

- application container(s)
- Storage resources
- Unique network IP

Shared Resources:

- Networking
- Storage



Example

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   creationTimestamp: null
5   labels:
6     run: nginx
7   name: nginx
8 spec:
9   containers:
10    - image: nginx
11      name: nginx
12      resources: {}
13    dnsPolicy: ClusterFirst
14    restartPolicy: Always
15 status: {}
```



Example

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   labels:
5     run: pod-busybox
6   name: pod-busybox
7 spec:
8   containers:
9     - command:
10       - sh
11       - -c
12         echo App is running! && sleep 30
13     image: busybox
14     name: pod-busybox
15     resources: {}
16 restartPolicy: Never
```



POD - Lifecycle

- Phase
 - Pending (waiting to be scheduled, image downloading)
 - Running (all containers started and ready to serve)
 - Succeeded (all containers exited with success)
 - Failed (all containers exited but at least one with failure)
 - Unknown (unable to fetch status as node is unreachable)
- Container States
 - Waiting, Running, Terminated
- Restart Policy (**Always**, Never, OnFailure)
- Conditions
 - Type: PodScheduled, ContainersReady, Initialized, Ready
(lastProbeTime, lastTransitionTime, Message, reason, status)
- Probes
 - Startup, Readiness, Liveness
- Lifecycle hooks



Phase - Pending

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   creationTimestamp: null
5   labels:
6     run: nginx
7     name: nginx
8 spec:
9   containers:
10  - image: nginx
11    name: nginx
12    resources:
13      requests:
14        cpu: "1000m"
15        memory: "1Gi"
16    dnsPolicy: ClusterFirst
17    restartPolicy: Never
18 status: {}
```



Phase - Running

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   creationTimestamp: null
5   labels:
6     run: busybox
7   name: busybox
8 spec:
9   containers:
10  - command:
11    - ping
12    - google.com
13    image: busybox
14    name: busybox
15    resources: {}
16  dnsPolicy: ClusterFirst
17  restartPolicy: Always
18 status: {}
```



Phase - Succeeded

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   creationTimestamp: null
5   labels:
6     run: busybox
7   name: busybox
8 spec:
9   containers:
10  - image: busybox
11    name: busybox
12    resources: {}
13  dnsPolicy: ClusterFirst
14  restartPolicy: Never
15 status: {}
```



Probes

- Types
 - Startup
 - Readiness
 - Liveness
- Methods
 - Http
 - Tcp
 - Command
- Settings
 - initialDelaySeconds
 - periodSeconds
 - timeoutSeconds
 - successThreshold
 - failureThreshold
- Http
 - Host
 - Scheme
 - Path
 - Port
 - Headers



Probe - Liveness - Exec

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: probe-liveness-exec
5 spec:
6   containers:
7     - name: probe-liveness-exec
8       image: k8s.gcr.io/busybox
9       args:
10        - /bin/sh
11        - -c
12        - touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep 600
13       livenessProbe:
14         exec:
15           command:
16            - cat
17            - /tmp/healthy
18         initialDelaySeconds: 5
19         periodSeconds: 5
```




Probe - Liveness - http

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: probe-liveness-http
5 spec:
6   containers:
7     - name: probe-liveness-http
8       image: k8s.gcr.io/liveness
9       args:
10        - /server
11       livenessProbe:
12         httpGet:
13           path: /healthz
14           port: 8080
15           httpHeaders:
16             - name: Custom-Header
17               value: Awesome
18         initialDelaySeconds: 3
19         periodSeconds: 3
```



Probe - Liveness - readiness - http

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: probe-liveness-readiness-tcp
5 spec:
6   containers:
7     - name: probe-liveness-readiness-tcp
8       image: k8s.gcr.io/goproxy:0.1
9       ports:
10        - containerPort: 8080
11       readinessProbe:
12         tcpSocket:
13           port: 8080
14         initialDelaySeconds: 5
15         periodSeconds: 10
16       livenessProbe:
17         tcpSocket:
18           port: 8080
19         initialDelaySeconds: 15
20         periodSeconds: 20
```



Probe - Liveness - startup - http

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: probe-liveness-startup-http
5 spec:
6   containers:
7     - name: probe-liveness-startup-http
8       image: k8s.gcr.io/liveness
9       args:
10        - /server
11       livenessProbe:
12         httpGet:
13           path: /healthz
14           port: 8080
15         failureThreshold: 1
16         periodSeconds: 10
17       startupProbe:
18         httpGet:
19           path: /healthz
20           port: 8080
21         failureThreshold: 30
22         periodSeconds: 10
```



POD Init Containers

- Always run to completion
- Must complete successfully before next one
- Readiness probes not supported
- Run(s) before application containers

Examples:

- Custom code / utilities to run before app containers
- Block / delay app container startup
- App container image building can be separate



POD Init - Statuses

- Init:N/M
- Init:Error
- Init:CrashLoopBackOff
- Pending
- PodInitializing
- Running



Example

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: init-containers
5 spec:
6   containers:
7     - name: main-container
8       image: busybox:1.28
9       command: ['sh', '-c', 'echo The app is running! && sleep 3600']
10  initContainers:
11    - name: init-service
12      image: busybox:1.28
13      command: ['sh', '-c', "until nslookup myservice.$(cat /var/run/secrets/kubernetes.io/serviceaccount/namespace).
svc.cluster.local; do echo waiting for myservice; sleep 2; done"]
14    - name: init-mydb
15      image: busybox:1.28
16      command: ['sh', '-c', "until nslookup mydb.$(cat /var/run/secrets/kubernetes.io/serviceaccount/namespace).svc.c
luster.local; do echo waiting for mydb; sleep 2; done"]
```




POD - Topology Spread

- Objectives
 - To control how Pods are spread across **regions**, zones, nodes and other user defined topology domains
 - To achieve high availability
 - To achieve efficient resource utilization
- Spread Constraints
 - `maxSkew`
 - `topologyKey`
 - `whenUnsatisfiable` (`DoNotSchedule` / `ScheduleAnyway`)
 - `labelSelector`



POD - Topology Spread






POD - Topology Spread

```
1 kind: Pod
2 apiVersion: v1
3 metadata:
4   name: topology-zone-constraint-1
5   labels:
6     foo: bar
7 spec:
8   topologySpreadConstraints:
9     - maxSkew: 1
10     topologyKey: failure-domain.beta.kubernetes.io/zone
11     whenUnsatisfiable: DoNotSchedule
12     labelSelector:
13       matchLabels:
14         foo: bar
15   containers:
16     - name: pause
17       image: k8s.gcr.io/pause:3.1
```

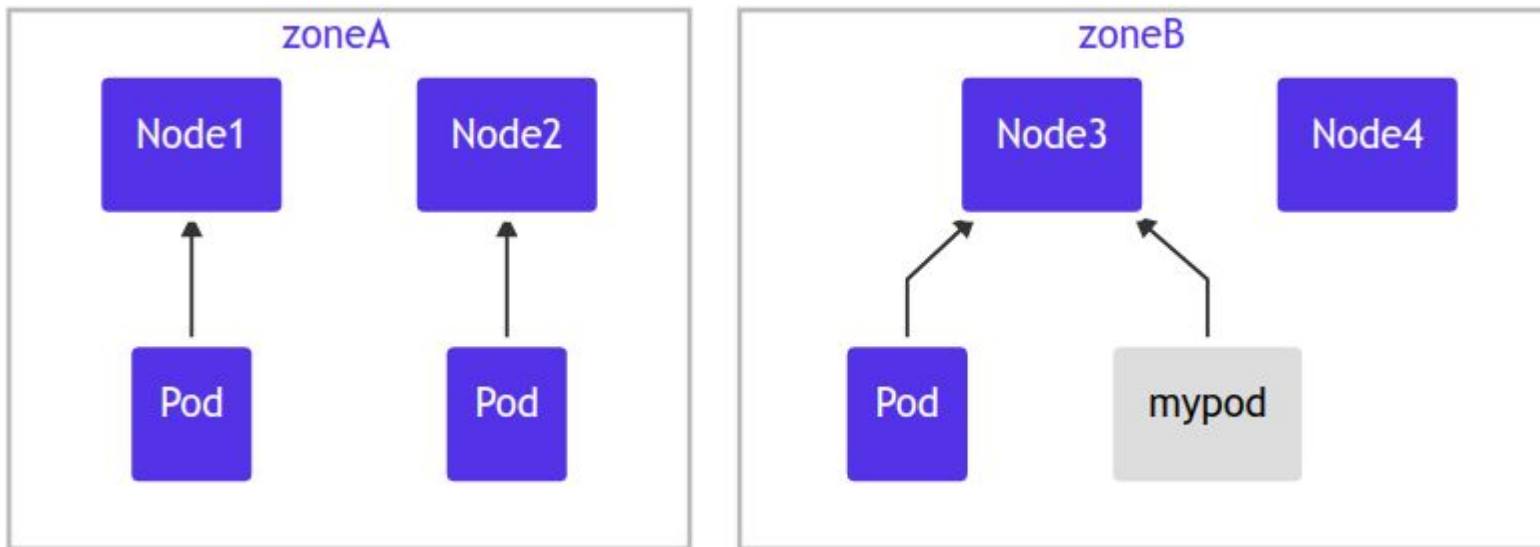


POD - Topology Spread





POD - Topology Spread





Topology - Multiple Constraints

```
1 kind: Pod
2 apiVersion: v1
3 metadata:
4   name: topology-constraints-two
5   labels:
6     foo: bar
7 spec:
8   topologySpreadConstraints:
9     - maxSkew: 1
10     topologyKey: topology.kubernetes.io/zone
11     whenUnsatisfiable: DoNotSchedule
12     labelSelector:
13       matchLabels:
14         foo: bar
15     - maxSkew: 1
16     topologyKey: kubernetes.io/hostname
17     whenUnsatisfiable: DoNotSchedule
18     labelSelector:
19       matchLabels:
20         foo: bar
21   containers:
22     - name: pause
23       image: k8s.gcr.io/pause:3.1
```



Topology - Constraint & Node Affinity

```
1 kind: Pod
2 apiVersion: v1
3 metadata:
4   name: topology-constraint-nodeaffinity
5   labels:
6     foo: bar
7 spec:
8   topologySpreadConstraints:
9     - maxSkew: 1
10     topologyKey: zone
11     whenUnsatisfiable: DoNotSchedule
12     labelSelector:
13       matchLabels:
14         foo: bar
15   affinity:
16     nodeAffinity:
17       requiredDuringSchedulingIgnoredDuringExecution:
18         nodeSelectorTerms:
19           - matchExpressions:
20             - key: zone
21               operator: NotIn
22               values:
23                 - zoneC
24   containers:
25     - name: pause
26       image: k8s.gcr.io/pause:3.1
```



POD Preset

- To inject certain info at creation time
- Can include secrets, volume, volume mounts, environment variables

```
1 apiVersion: settings.k8s.io/v1alpha1
2 kind: PodPreset
3 metadata:
4   name: pod-present-db-config
5 spec:
6   selector:
7     matchLabels:
8       role: frontend
9   env:
10    - name: DB_PORT
11      value: "6379"
12    - name: duplicate_key
13      value: FROM_ENV
14    - name: expansion
15      value: $(REPLACE_ME)
16   envFrom:
17    - configMapRef:
18      name: etcd-env-config
19   volumeMounts:
20    - mountPath: /cache
21      name: cache-volume
22   volumes:
23    - name: cache-volume
24      emptyDir: {}
25
```

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: pod-preset-config
5 data:
6   number_of_members: "1"
7   initial_cluster_state: new
8   initial_cluster_token: DUMMY_ETCD_INITIAL_CLUSTER_TOKEN
9   discovery_token: DUMMY_ETCD_DISCOVERY_TOKEN
10  discovery_url: http://etcd_discovery:2379
11  etcdctl_peers: http://etcd:2379
12  duplicate_key: FROM_CONFIG_MAP
13  REPLACE_ME: "a value"
```



POD - Ephemeral

- Meant for interactive troubleshooting inside POD
- No resource guarantees
- Never restart automatically
- Process Namespace sharing

```
kubectl replace --raw /api/v1/namespaces/default/pods/example-pod/ephemeralcontainers -f ec.json
```

```
{
  "apiVersion": "v1",
  "kind": "EphemeralContainers",
  "metadata": {
    "name": "example-pod"
  },
  "ephemeralContainers": [{
    "command": [
      "sh"
    ],
    "image": "busybox",
    "imagePullPolicy": "IfNotPresent",
    "name": "debugger",
    "stdin": true,
    "tty": true,
    "terminationMessagePolicy": "File"
  }]
}
```



Exercises (30 mins)

- Create POD with init containers
 - Main app container from tutum/hello-world
 - Init container using busybox to fetch `google.com/index.html` and save to `/www/google.html`
- Create POD to keep in different state (Pending, Running, Succeeded)
- Create POD with readiness probe for container readiness requiring 60 seconds
- Create POD and restart if not live after 30 seconds having 10 s interval
- Create POD that requires 90 s to start but restart every 30 s if not live
- Create two PODs and make sure they are not deployed on the same node
- Create two PODs and make sure they are deployed on the same node
- Create two PODs and ensure they are deployed on different zones
- Create four PODs and ensure no more than one is deployed on same zone



Multi Container Pods

- Ambassador
- Sidecar
- Adapter
- Demo
- Practicals



Patterns - POD

- To extend the functionality of the existing container
- To have helper process enhancing work of the existing container
- To send logs to external server
- Types
 - Sidecar - To export logs
 - Ambassador - To proxy connection
 - Adapter - To standardise and normalize output



Pattern - Sidecar

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: sidecar
5 spec:
6   volumes:
7     - name: shared-logs
8       emptyDir: {}
9   containers:
10    - name: main-container
11      image: alpine
12      command: ["/bin/sh"]
13      args: ["-c", "while true; do date >> /var/log/index.html; sleep 10;done"]
14      volumeMounts:
15        - name: shared-logs
16          mountPath: /var/log
17    - name: sidecar-container
18      image: nginx
19      ports:
20        - containerPort: 80
21      volumeMounts:
22        - name: shared-logs
23          mountPath: /usr/share/nginx/html
```



Pattern - Adapter

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: adapter
5 spec:
6   volumes:
7     - name: shared-logs
8       emptyDir: {}
9   containers:
10    - name: main-container
11      image: alpine
12      command: ["/bin/sh"]
13      args: ["-c", "while true; do date > /var/log/top.txt && top -n 1 -b >> /var/log/top.txt; sleep 10;done"]
14      volumeMounts:
15        - name: shared-logs
16          mountPath: /var/log
17    - name: adapter-container
18      image: alpine
19      command: ["/bin/sh"]
20      args: ["-c", "while true; do (cat /var/log/top.txt | head -1 > /var/log/status.txt) && (cat /var/log/top.txt | head -2 | tail -1 | grep
21 -o -E '\\d+\\w' | head -1 >> /var/log/status.txt) && (cat /var/log/top.txt | head -3 | tail -1 | grep
22 -o -E '\\d+%' | head -1 >> /var/log/status.txt); sleep 5; done"]
23      volumeMounts:
24        - name: shared-logs
25          mountPath: /var/log
```



Pattern - Ambassador

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: ambassador-nginx-config
5 data:
6   nginx.conf: |
7     worker_processes 1;
8     worker_rlimit_nofile 4096;
9     events {
10       worker_connections 512;
11     }
12     http {
13       proxy_set_header Host $host;
14       proxy_set_header X-Forwarded-Proto $scheme;
15       proxy_set_header X-Real-IP $remote_addr;
16       proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
17       upstream backend {
18         server msn.com:80;
19       }
20       server {
21         listen 80;
22         location / {
23           proxy_pass http://backend;
24         }
25       }
26     }
27 ---
28 apiVersion: v1
29 kind: Pod
30 metadata:
31   name: multi-pod-ambassador
32 spec:
33   containers:
34     - name: main-app
35       image: busybox
36       imagePullPolicy: IfNotPresent
37       command: ["/bin/sh"]
38       args: ["-c", "while true;do wget -O /tmp/app.txt localhost ;sleep 30;done"]
39     - name: ambassador
40       image: nginx
41       imagePullPolicy: IfNotPresent
42       ports:
43         - containerPort: 80
44       volumeMounts:
45         - name: nginx-config
46           mountPath: /etc/nginx/nginx.conf
47           subPath: nginx.conf
48   volumes:
49     - name: nginx-config
50       configMap:
51         name: ambassador-nginx-config
```



Exercises

- Create POD having below
 - Main container: busybox
 - Appending ping outputs to /tmp/index-input.html
 - Sidecar container: busybox
 - Copy the index.html to /tmp/index-output.html
- Create POD having below
 - Main container: alpine
 - Print date to /tmp/index.html
 - Adapter container: busybox
 - Copy the index.html to /tmp/index-output.html
- Create POD having below
 - Main container: busybox
 - Storing localhost ping outputs to /tmp/index-input.html
 - Ambassador container: nginx
 - Proxying localhost:80 to <https://www.brainupgrade.in>



Pod Design

- Deployments
- Rolling Updates & Rollbacks
- App Deployment - multi node
- Auto scaling pods
- Auto scaling pod across availability zones
- Jobs, Cron Jobs
- Labels, Selectors, Annotations
- Demo
- Practicals



Deployments

Use Cases

- To rollout a set of PODs
- To declare a new set of PODs
- To rollback to an earlier version of deployment
- To scale up deployment to facilitate more load
- To pause the deployment / rollout
- To autoscale deployment when cpu usage threshold reached



Deployment - Example

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: test-app
5 spec:
6   replicas: 1
7   selector:
8     matchLabels:
9       component: test-app
10  template:
11    metadata:
12      labels:
13        component: test-app
14    spec:
15      containers:
16        - name: test-app
17          image: brainupgrade/test-app:all-tiers-in-one
18          imagePullPolicy: IfNotPresent
19          ports:
20            - containerPort: 8080
21      resources:
22        requests:
23          cpu: "100m"
24          memory: "250Mi"
25
```



Deployment - Commands

- `kubectl create deployment nginx --image=nginx:1.15 --replicas=5`
- `kubectl get deployment/nginx`
- `kubectl describe deployment/nginx`
- `kubectl rollout history deployment/nginx`
- `kubectl set image deployment/nginx nginx=nginx:1.16`
- `kubectl rollout history deployment/nginx`
- `kubectl rollout undo deployment/nginx`
- `kubectl rollout undo deployment/nginx --to-revision=2`
- `kubectl scale --replicas=50 deployment/nginx`
- `kubectl rollout pause deployment/nginx`
- `kubectl rollout status deployment/nginx`
- `kubectl rollout resume deployment/nginx`
- `kubectl autoscale deployment/nginx --min=2 --max=10`



Job

- To provide reliable parallel execution of tasks
- Examples:
 - Send emails, transcode files, Scan database for a set of rows,
- Patterns
 - Non parallel job
 - Fixed completion count job
 - Work queue job



Job

```
1 apiVersion: batch/v1
2 kind: Job
3 metadata:
4   name: job
5 spec:
6   template:
7     spec:
8       containers:
9         - name: perl
10           image: perl
11           command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
12       restartPolicy: Never
13   backoffLimit: 4
```



Job - Timeout

```
1 apiVersion: batch/v1
2 kind: Job
3 metadata:
4   name: job-timeout
5 spec:
6   backoffLimit: 5
7   activeDeadlineSeconds: 100
8   template:
9     spec:
10      containers:
11      - name: perl
12        image: perl
13        command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
14      restartPolicy: Never
```



CronJob

- Creates jobs on a repeating schedule
- Schedule times are based on kube-controller-manager
- Useful for tasks like migrating data to reporting server, sending emails, creating backups etc
- Schedule tasks at specific time (like when cluster is idle)

Key Configurations:

- startingDeadlineSeconds - Missed occurrences in last X seconds will be counted
- concurrencyPolicy
 - If Allow, then job will run at least once
 - If Forbid, will be missed if previous instance is still running



CronJob - Expression

```
# |_____ minute (0 - 59)
# | |_____ hour (0 - 23)
# | | |_____ day of the month (1 - 31)
# | | | |_____ month (1 - 12)
# | | | | |_____ day of the week (0 - 6) (Sunday to Saturday;
# | | | | | 7 is also Sunday on some systems)
# | | | | |
# | | | | |
# | | | | |
# | * * * * * <command to execute>
```

Examples:

- `*/15 0,8,16 * * * echo running backup` (every 15 minutes of 0,8 & 16th hour)
- `30 0 * * 6 /home/oracle/scripts/export_dump.sh` (last day of week at 00:30)
- `1 0 * * * printf "" > /var/log/apache/error_log` (everyday at 00:01)



CronJob - Expression

```
1 apiVersion: batch/v1beta1
2 kind: CronJob
3 metadata:
4   name: cron-job
5 spec:
6   schedule: "*/1 * * * *"
7   jobTemplate:
8     spec:
9       template:
10        spec:
11          containers:
12            - name: cron-job
13              image: busybox
14              args:
15                - /bin/sh
16                - -c
17                - date; echo Migrating data to reporting server...
18          restartPolicy: OnFailure
19
```




Labels

- Labels
 - Key value pairs attached to objects
 - To specify identifying attributes of objects
 - To organize and select subset of objects
 - To query objects efficiently (cli as well gui monitoring tools)
 - Attached at creation time and can be added / modified at any time
 - Label key must be unique per object
- Example labels:
 - `"release" : "stable", "release" : "canary"`
 - `"environment" : "dev", "environment" : "qa", "environment" : "production"`
 - `"tier" : "frontend", "tier" : "backend", "tier" : "cache"`
 - `"partition" : "customerA", "partition" : "customerB"`
 - `"track" : "daily", "track" : "weekly"`



Labels

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-labels
5   labels:
6     environment: production
7     app: nginx
8 spec:
9   containers:
10    - name: nginx
11      image: nginx
12      ports:
13        - containerPort: 80
```



Selectors

- Equality Based

- `environment = production`
- `tier != frontend`

- Set Based

- `environment in (production, qa)`
- `tier notin (frontend, backend)`
- `partition`
- `!partition`



Selectors - Examples

- `kubectl get pods -l environment=production,tier=frontend`
- `kubectl get pods -l 'environment in (production),tier in (frontend)'`
- `kubectl get pods -l 'environment in (production, qa)'`
- `kubectl get pods -l 'environment,environment notin (frontend)'`

Jobs, Deployments, ReplicaSet, Daemonset

```
selector:  
  matchLabels:  
    component: redis  
  matchExpressions:  
    - {key: tier, operator: In, values: [cache]}  
    - {key: environment, operator: NotIn, values: [dev]}
```



Selectors - Examples

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: selector-pod-node
5 spec:
6   containers:
7     - name: cuda-test
8       image: "k8s.gcr.io/cuda-vector-add:v0.1"
9       resources:
10        limits:
11          nvidia.com/gpu: 1
12   nodeSelector:
13     accelerator: nvidia-tesla-p100
```



Annotations

- To attach non-identifying arbitrary metadata to objects
- Usage
 - Pointers for debugging purposes
 - Build, release, image hashes etc
 - Author info, contact details
 - Metadata to help tools for deployment, management, introspection

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: annotations-pod
5   annotations:
6     imageregistry: "https://hub.docker.com/"
7 spec:
8   containers:
9     - name: nginx
10      image: nginx
11      ports:
12        - containerPort: 80
```



Exercises (30 mins)

- Create a deployment to run 3 replicas of nginx container
- Scale down the replicas to 1
- Scale up replicas to 10
- View the roll out history
- Switch to rollout version 2
- View deployments, rc, pod using kubectl
- Scale down replicas to 5
- Update image to nginx:1.18 and immediately try another rollout with nginx:1.17
- Observe if rollout with nginx:1.17 was completed or not
- Autoscale pod to max 5 min 1



Persistence

- Persistence Volume
- Persistence Volume Claim
- Statefulset
- Daemonset



Overview - PV / PVC

- Ephemeral
 - Tightly coupled with POD lifetime
 - Deleted when POD is removed
 - Example: emptydir
- Persistent
 - Survives POD reboots
 - Meant for long term and independent of POD / Node lifecycle
 - Examples: hostpath, NFS, Cloud storage (EBS etc)
- The access modes are:
 - ReadWriteOnce -- the volume can be mounted as read-write by a single node
 - ReadOnlyMany -- the volume can be mounted read-only by many nodes
 - ReadWriteMany -- the volume can be mounted as read-write by many nodes



Examples - emptyDir

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: volume-emptydir
5 spec:
6   containers:
7     - image: nginx
8       name: test-container
9       volumeMounts:
10      - mountPath: /cache
11        name: cache-volume
12   volumes:
13     - name: cache-volume
14       emptyDir: {}
```



Examples - hostpath (file/dir)

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: hostpath-volume
5 spec:
6   containers:
7   - image: nginx
8     name: test-container
9     volumeMounts:
10    - mountPath: /data-mounted-as
11      name: hostpath-volume
12   volumes:
13   - name: hostpath-volume
14     hostPath:
15       # directory location on host
16       path: /data
17       # this field is optional
18       type: DirectoryOrCreate
```



Persistent Volume - local & node


```
1 apiVersion: v1
2 kind: PersistentVolume
3 metadata:
4   name: pv-local-node-affinity
5 spec:
6   capacity:
7     storage: 10Gi
8   # volumeMode field requires BlockVolume Alpha feature gate to be enabled.
9   volumeMode: Filesystem
10  accessModes:
11    - ReadWriteMany
12  persistentVolumeReclaimPolicy: Delete
13  storageClassName: local-storage
14  local:
15    path: /mnt/disks/ssd1
16  nodeAffinity:
17    required:
18      nodeSelectorTerms:
19        - matchExpressions:
20          - key: kubernetes.io/hostname
21            operator: In
22            values:
23              - ip-172-31-87-231.ec2.internal
```



Persistent Volume - EBS

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: volume-ebs
5 spec:
6   containers:
7     - image: k8s.gcr.io/test-webserver
8       name: test-container
9       volumeMounts:
10        - mountPath: /test-ebs
11          name: ebs-volume
12   volumes:
13     - name: ebs-volume
14       # This AWS EBS volume must already exist.
15       awsElasticBlockStore:
16         volumeID: <volume-id>
17         fsType: ext4
```

Persistent Volume - static






Example

```
1 apiVersion: v1
2 kind: PersistentVolume
3 metadata:
4   name: pv-hostpath
5   annotations:
6     pv.beta.kubernetes.io/gid: "1234"
7   labels:
8     type: local
9 spec:
10  storageClassName: manual
11  capacity:
12    storage: 10Gi
13  accessModes:
14    - ReadWriteOnce
15  hostPath:
16    path: "/mnt/data"
```

```
1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4   name: pvc-test
5 spec:
6   storageClassName: manual
7   accessModes:
8     - ReadWriteOnce
9   resources:
10     requests:
11       storage: 3Gi
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-pv-pvc
5 spec:
6   volumes:
7     - name: pv-storage
8       persistentVolumeClaim:
9         claimName: pvc-test
10  containers:
11    - name: task-pv-container
12      image: nginx
13      ports:
14        - containerPort: 80
15          name: "http-server"
16      volumeMounts:
17        - mountPath: "/usr/share/nginx/html"
18          name: pv-storage
```

Persistent Volume - Dynamic





Example

```
1 kind: StorageClass
2 apiVersion: storage.k8s.io/v1
3 metadata:
4   name: storageclass-generic
5 provisioner: kubernetes.io/aws-ebs
6 parameters:
7   type: gp2
8   zones: us-east-1a,us-east-1b,us-east-1c
9   iopsPerGB: "10"
10  fsType: ext4
```

```
1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4   name: pvc-dynamic
5   labels:
6     app: nginx
7 spec:
8   storageClassName: storageclass-generic
9   accessModes:
10    - ReadWriteOnce
11   resources:
12     requests:
13       storage: 1Gi
```

```
1 kind: Pod
2 apiVersion: v1
3 metadata:
4   name: pod-volume-dynamic
5   labels:
6     app: nginx
7 spec:
8   containers:
9     - name: nginx
10       image: nginx
11       volumeMounts:
12         - mountPath: "/var/www/html"
13           name: external
14   volumes:
15     - name: external
16       persistentVolumeClaim:
17         claimName: pvc-dynamic
```



Persistent Volumes

- GCEPersistentDisk
- AWSElasticBlockStore
- AzureFile
- AzureDisk
- CSI
- FC (Fibre Channel)
- FlexVolume
- Flocker
- NFS
- iSCSI
- RBD (Ceph Block Device)
- CephFS
- Cinder (OpenStack block storage)
- Glusterfs
- VsphereVolume
- Quobyte Volumes
- HostPath (Single node testing only – local storage is not supported in any way and WILL NOT WORK in a multi-node cluster)
- Portworx Volumes
- ScaleIO Volumes
- StorageOS



StatefulSet

Use Cases

- Stable, unique network identifiers
- Stable, persistent storage
- Ordered, graceful deployment and scaling
- Ordered, automated rolling updates

Limitations

- Requires headless service (manual way)
- No automatic deletion of referenced volumes
- No PODs deletion guarantee when StatefulSet is deleted
- Rolling Updates not consistent always

StatefulSet - Example

```
1 apiVersion: apps/v1
2 kind: StatefulSet
3 metadata:
4   name: sts-web
5 spec:
6   serviceName: "nginx"
7   replicas: 2
8   selector:
9     matchLabels:
10      app: nginx
11   template:
12     metadata:
13       labels:
14         app: nginx
15     spec:
16       containers:
17         - name: nginx
18           image: k8s.gcr.io/nginx-slim:0.8
19           ports:
20             - containerPort: 80
21               name: web
22           volumeMounts:
23             - name: www
24               mountPath: /usr/share/nginx/html
25   volumeClaimTemplates:
26     - metadata:
27       name: www
28     spec:
29       accessModes: [ "ReadWriteOnce" ]
30       resources:
31         requests:
32           storage: 1Gi
```

- Scale Up
- Scale Up
- Update (image)



StatefulSet - Example

- Scale Up
- Scale down
- `kubectl set image sts/sts-web nginx=nginx:1.18`
- Staged Update
 - `kubectl patch statefulset sts-web -p '{"spec":{"updateStrategy":{"type":"RollingUpdate","rollingUpdate":{"partition":3}}}}'`
 - `kubectl patch statefulset sts-web --type=json -p='[{"op": "replace", "path": "/spec/template/spec/containers/0/image", "value":"nginx:1.17"}]'`
 - `kubectl get pod sts-web-1 --template '{{range $i, $c := .spec.containers}}{{ $c.image }}{{end}}'`



Controllers - DaemonSet

Purpose

- To run a copy of a POD on all / some node(s)

Use Cases

- Storage cluster daemon (gluster, ceph)
- Log Collectors (fluentd, logstash)
- Node Monitoring daemons (Prometheus, Dynatrace, collectd)



DaemonSet - Example

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd-elasticsearch
  namespace: kube-system
  labels:
    k8s-app: fluentd-logging
spec:
  selector:
    matchLabels:
      name: fluentd-elasticsearch
  template:
    metadata:
      labels:
        name: fluentd-elasticsearch
    spec:
      tolerations:
        - key: node-role.kubernetes.io/master
          effect: NoSchedule
      containers:
        - name: fluentd-elasticsearch
          image: quay.io/fluentd_elasticsearch/fluentd:v2.5.2
          resources:
            limits:
              memory: 200Mi
            requests:
              cpu: 100m
              memory: 200Mi
          volumeMounts:
            - name: varlog
              mountPath: /var/log
            - name: varlibdockercontainers
              mountPath: /var/lib/docker/containers
              readOnly: true
          terminationGracePeriodSeconds: 30
      volumes:
        - name: varlog
          hostPath:
            path: /var/log
        - name: varlibdockercontainers
          hostPath:
            path: /var/lib/docker/containers
```



Exercises (15 mins)

- Create POD (nginx / redis) to use volume emptyDir
- Launch POD and login into POD
- Create test file
- Kill the container process (nginx / redis)
- Observe POD status and login into POD again
- Verify if test file exists

NOTE:

POD has restartPolicy as Always

Ephemeral storage is associated till POD is deleted



Exercises (15 mins)

- Create nginx POD that uses pvc for serving web files
- Define pvc that uses pv
- Define pv that refers to host path /mnt/data
- Create index.html echoing 'hello k8s' under host path
- Verify that nginx serves the index.html contents that you saved



Exercises (15 mins)

- Create StatefulSet having image=nginx:1.16 and replicas 5
- Scale replicas to 5
- Change image to nginx:1.17
- Set the rolling update partition to 3
- Change the image to nginx:1.17
- Watch the pods scaling up / down using
 - `kubectl get -w pods`



Configuration

- Config Maps, Environment
- Secrets
- Security Contexts
- Service Accounts
- Demo
- Practicals

ConfigMap

- To store non-confidential key-value pairs
- Can be consumed as env variables, command line args or config files in volume
- To decouple env specific config from images for portability
- Max data 1MB

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: cm-game-demo
5 data:
6   # property-like keys; each key maps to a simple value
7   player_initial_lives: "3"
8   ui_properties_file_name: "user-interface.properties"
9   # file-like keys
10  game.properties: |
11    enemy.types=aliens,monsters
12    player.maximum-lives=5
13  user-interface.properties: |
14    color.good=purple
15    color.bad=yellow
16    allow.textmode=true
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-configmap
5 spec:
6   containers:
7     - name: demo
8       image: alpine
9       command: ["sleep", "3600"]
10      env:
11        # Define the environment variable
12        - name: PLAYER_INITIAL_LIVES # Notice that the case is different here
13          valueFrom: # from the key name in the ConfigMap.
14            configMapKeyRef:
15              name: cm-game-demo # The ConfigMap this value comes from.
16              key: player_initial_lives # The key to fetch.
17        - name: UI_PROPERTIES_FILE_NAME
18          valueFrom:
19            configMapKeyRef:
20              name: cm-game-demo
21              key: ui_properties_file_name
22      volumeMounts:
23        - name: config
24          mountPath: "/config"
25          readOnly: true
26      volumes:
27        # You set volumes at the Pod level, then mount them into containers inside that Pod
28        - name: config
29          configMap:
30            # Provide the name of the ConfigMap you want to mount.
31            name: cm-game-demo
32            # An array of keys from the ConfigMap to create as files
33            items:
34              - key: "game.properties"
35                path: "game.properties"
36              - key: "user-interface.properties"
37                path: "user-interface.properties"
```



Secret

- To manage sensitive info like password, oauthkeys, docker login, ssh keys, tls etc

- Examples

```
kubectl create secret docker-registry secret-tiger-docker \  
  --docker-username=tiger \  
  --docker-password=pass113 \  
  --docker-email=tiger@acme.com
```

```
kubectl create secret tls my-tls-secret \  
  --cert=path/to/cert/file \  
  --key=path/to/key/file
```



Secret

```
1 apiVersion: v1
2 data:
3   username: YWRtaW4=
4   password: MWYyZDFlMmU2N2Rm
5 kind: Secret
6 metadata:
7   name: pod-secret
8   namespace: default
9   resourceVersion: "164619"
10  uid: cfee02d6-c137-11e5-8d73-42010af00002
11 type: Opaque
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-secret
5 spec:
6   containers:
7     - name: mypod
8       image: redis
9       env:
10        - name: SECRET_USERNAME
11          valueFrom:
12            secretKeyRef:
13              name: mysecret
14              key: username
15        - name: SECRET_PASSWORD
16          valueFrom:
17            secretKeyRef:
18              name: mysecret
19              key: password
20       volumeMounts:
21        - name: foo
22          mountPath: "/etc/foo"
23          readOnly: true
24       volumes:
25        - name: foo
26          secret:
27            secretName: pod-secret
28            items:
29              - key: username
30                path: my-group/my-username
```



Security Context

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: security-context-demo
5 spec:
6   securityContext:
7     runAsUser: 1000
8     runAsGroup: 3000
9     fsGroup: 2000
10  volumes:
11    - name: sec-ctx-vol
12      emptyDir: {}
13  containers:
14    - name: sec-ctx-demo
15      image: busybox
16      command: [ "sh", "-c", "sleep 1h" ]
17      volumeMounts:
18        - name: sec-ctx-vol
19          mountPath: /data/demo
20      securityContext:
21        allowPrivilegeEscalation: false
22        capabilities:
23          add: ["NET_ADMIN", "SYS_TIME"]
24
```



User Accounts & Service Accounts

- User Accounts
 - User accounts are for humans.
 - User accounts are intended to be global. Names must be unique across all namespaces of a cluster.
- Service Accounts
 - Service accounts are for processes, which run in pods.
 - Service accounts are namespaced.
 - Service account creation is intended to be more lightweight



Scheduling

- Node Name
- Taints
- Tolerations
- Affinity
- Demo
- Practicals



Topology

Purpose

- To route the traffic in the same node / cluster / zone wherever possible

Advantages

- Lower latency
- Cost Optimization as inter zonal cloud requests costs

How to achieve

- Enable Service Topology feature (API Server and Kube proxy)
- Label endpoints, node, cluster, zone appropriately



Scheduling - Priorities

- Node capacity based on priority weightage on its resources (cpu, memory, disk etc)
- Node score - Input to Kube scheduler
- Resource specification on PODs helps better computing capacity management



Services

- Cluster IP
- Node Port
- Load Balancer
- Connecting using services
- Demo
- Practicals



Service

- An abstract way to expose an application running on pod as network service.
- Frontends and backends of application can connect without worrying about POD IPs

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```



Service without selector

Use Cases:

- External DB Cluster in production
- To point a service in another namespace / cluster
- Uses session affinity while connecting to backend PODs

```
apiVersion: v1
kind: Endpoints
metadata:
  name: my-service
subsets:
  - addresses:
    - ip: 192.0.2.42
    ports:
    - port: 9376
```



Service Types

- Cluster IP
 - Service exposed on cluster internal IP
 - Reachable only within cluster
- Node Port
 - Exposed on each Node IP at static port
- Load Balancer
 - Exposed through external cloud load balancer
- External Name
 - Exposed through the contents of external field via CNAME record



External IP

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - name: http
      protocol: TCP
      port: 80
      targetPort: 9376
  externalIPs:
    - 80.11.12.10
```




Ingress & Networking

- Ingress Controller
- Ingress Resources
- Network Policies
- Host mapping with service
- TLS
- Multi domain frontends with service
- Demo
- Practicals



DNS

DNS Policy

- Default
 - Inherit name resolution from Node
- ClusterFirst
 - Forwards to upstream nameserver for unresolved name queries
- ClusterFirstWithHostNet
 - Only for PODs running with hostNetwork
- None
 - POD explicitly defines it using dnsConfig



Custom DNS - Example

```
apiVersion: v1
kind: Pod
metadata:
  namespace: default
  name: dns-example
spec:
  containers:
    - name: test
      image: nginx
  dnsPolicy: "None"
  dnsConfig:
    nameservers:
      - 1.2.3.4
    searches:
      - ns1.svc.cluster-domain.example
      - my.dns.search.suffix
    options:
      - name: ndots
        value: "2"
      - name: edns0
```



Ingress

- Provides load balancing, SSL Termination and Name based virtual hosting
- Provides externally reachable URLs to Services
- Used for HTTP / HTTPS protocols

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: test-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
      paths:
      - path: /testpath
        backend:
          serviceName: test
          servicePort: 80
```



Examples

- Web server (tutum)
- API Services offering from UniGPS



Exercises (15 mins)

- Create deployment based on nginx image with 3 replicas
- Create NodePort service to map to the PODs created by above deployment
- View service URL and access it using browser
- Create another service of type Ingress
- View service URL and Access the service outside cluster
- Create one more service of type ClusterIP
- View service URL and find a way to access it



Exercise - Scenario

Assume that based on your recently acquired K8S expertise, you are tasked by your firm to develop real time video based fleet monitoring service with below high level

Objectives:

- New video service should be independent of any other earlier services (/API) developed
- Deployment should be as easy as possible
- New service should be provided to end customers via `video.unigps.in`
- You are expected to use current k8s setup and extend on it

Outcome expected:

- Yaml based definitions of deployment, service and domain based routing and load balancing.
- Service deployment should have at least 5 instances of replica



Monitoring

- Liveness Probes
- Readiness Probes
- Container Logging
- Monitoring & Debugging
 - Debug - live debugging with Telepresence
 - Shell to running container
- Demo
- Practicals



Monitoring Dashboard

```
kubectl apply -f  
https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.0-beta8/aio/deploy/recommended.yaml
```

```
kubectl proxy
```

<http://localhost:8001/api/v1/namespaces/kubernetes-dashboard/services/https:kubernetes-dashboard:/proxy/>



Cluster Access

To view cluster configuration

```
kubectl config view
```

Reverse proxy to API server

```
kubectl proxy --port=8080
```



Port Forwarding

```
kubectl port-forward <pod> 7000:6379
```

```
kubectl port-forward <deployment> 7000:6379
```

```
kubectl port-forward <svc> 7000:6379
```

To access LoadBalancer service on localhost

```
minikube tunnel
```



Introspection & debugging

```
kubectl get pods
kubectl get pod <pod-name> -o yaml
kubectl describe <pod-name>
kubectl describe <pod-name> -o yaml
kubectl get events
kubectl get events --namespace=my-namespace (--all-namespaces)
kubectl get nodes
kubectl get node <node-name>
kubectl get node <node-name> -o yaml
kubectl describe node <node-name>
kubectl describe node <node-name> -o yaml
```

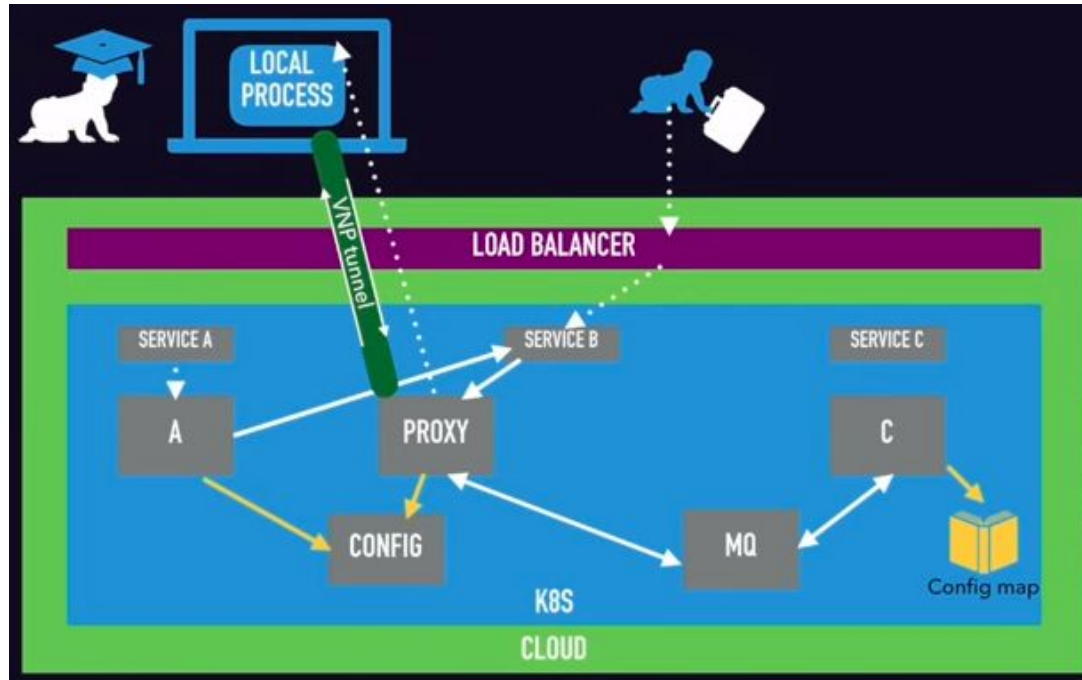


Live debugging using IDE

```
telepresence --swap-deployment hostnames --namespace default --run mvn spring-boot:run  
-Dspring-boot.run.jvmArguments="-Xdebug  
-Xrunjdwp:transport=dt_socket,server=y,suspend=y,address=5005"
```

Project: git/rest-service

Telepresence - Live debugging





Shell to running container

```
rajesh@rajesh-Gazelle:~/git/kubernetes/debugging/shell$ kubectl apply -f shell-demo.yaml
```

```
kubectl get pod shell-demo
```

```
kubectl exec -it shell-demo -- /bin/bash
```

```
root@shell-demo:/# ls /
```

```
root@shell-demo:/# echo Hello shell demo > /usr/share/nginx/html/index.html
```

```
root@shell-demo:/# apt-get update
```

```
root@shell-demo:/# apt-get install curl
```

```
root@shell-demo:/# curl localhost
```

```
kubectl exec shell-demo env
```

```
kubectl exec -it my-pod --container main-app -- /bin/bash
```



Best Practices

- Configuration - specify latest stable API version
- Keep config files in version control before pushing to cluster
- Prefer YAML over JSON
- Group related objects into one file whenever it makes more sense
- Don't specify default values unnecessarily
- Put Object descriptions as part of annotations
- Don't use naked PODs
- Create service before deployments
- Avoid using hostPort for POD
- Use labels effectively
- Use image tag instead of using latest as the default
- Use kubectl run and expose to launch single container deployments & services



Best Practices- Security

4 Cs of Cloud Native Security

- Code
- Container
- Cluster
- Cloud

Areas of caution / concern:

- API Server: Avoid exposing Master Nodes / API server publicly
- Nodes Access: Should allow only master nodes to communicate on specified ports
- Access to Cloud API: Based on principle of least privilege, Cloud API access to K8S should be provided
- ETCD: Master only should have access and data should be encrypted



Recap

- Review
- Q & A

Thank You for your active participation!

raiesh@unigps.in

9880195215

<https://www.linkedin.com/in/raiesh-g-b48495/>